

# Symbolic AI's for Generalised Connect Four

Bram Pulles – S1015194

October 7, 2022

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Evaluation Space</b>	<b>2</b>
2.1	Win, Loss, Draw . . . . .	2
2.2	Heuristic . . . . .	3
<b>3</b>	<b>Heuristics</b>	<b>3</b>
3.1	Given Heuristic . . . . .	3
3.2	Nop Heuristic . . . . .	4
<b>4</b>	<b>Negamax AI</b>	<b>4</b>
<b>5</b>	<b>Alpha-beta AI</b>	<b>5</b>
<b>6</b>	<b>Insane AI</b>	<b>5</b>
6.1	Move ordering . . . . .	5
6.2	Transposition table . . . . .	5
6.3	Iterative deepening . . . . .	7
6.4	MTD(f) . . . . .	7
<b>7</b>	<b>Complexity</b>	<b>10</b>
<b>8</b>	<b>Benchmarks</b>	<b>10</b>
8.1	Negamax AI . . . . .	10
8.2	Alpha-beta AI . . . . .	11
8.3	Insane AI . . . . .	11
<b>9</b>	<b>Conclusion</b>	<b>12</b>
<b>10</b>	<b>Discussion</b>	<b>12</b>
<b>11</b>	<b>References</b>	<b>12</b>
<b>A</b>	<b>Implementation</b>	<b>13</b>

# 1 Introduction


Connect Four is a widely known game, most of us probably played this as a kid for fun. What most people do not know is that Connect Four is actually strongly solved, i.e. for every state the best move is known [6]. In this assignment we are tasked with creating various search programs for Generalised Connect Four. In this version the board is of arbitrary width and height, and the number,  $N$ , of consecutive stones needed for a win is unspecified.

I have implemented a negamax AI, which is equivalent to minmax, an alpha-beta pruning AI and an insane AI which contains a lot of extra's such as move ordering, a transposition table, iterative deepening and MTD(f).

In this paper I start with an explanation of the evaluation space, as I made quite some nontrivial adjustments to the one given in the assignment. I then explain the negamax AI, the negamax AI with alpha-beta pruning and the further improved AI, pseudocode is used for all of these<sup>1</sup>. Next, I elaborate on the complexity of all of these algorithms and I discuss benchmark results for the different AI's. At last, I provide a conclusion and discussion of the results. In the appendix I have noted some small notes about my implementation.

## 2 Evaluation Space

Before we get into how all of the different algorithms work and how they are implemented we have to talk about evaluations. I have made a generalised evaluation space which is suited for the algorithms that I implemented.



The evaluation space is shown above. Every evaluation of a board state is represented by a single integer. This makes it extremely fast to compare different evaluations.

In the evaluation space we can see on the left side negative integers going to  $-\infty$ , while on the right side we have positive integers going to  $\infty$ . There are a lot of abbreviations in the diagram: L stands for Loss, D for Draw, W for Win, Hl for Heuristic loss and Hw for Heuristic win. We also have HL for Heuristic Limit, a heuristic value  $v$  always satisfies  $v \in [0, HL]$ , and ML for Movecount Limit which is the maximum number of moves that can be made, i.e. the width of the board times the height of the board.

### 2.1 Win, Loss, Draw

The evaluation space is carefully constructed such that desired outcomes have a higher value than less desired outcomes. The best possible situation is a win in zero moves, while the worst possible situation is a loss in zero moves.

Within the winning space we prefer evaluations which lead us to a win faster than other winning evaluations. The worst winning evaluation would be a win in ML moves.

Within the losing space we prefer evaluations which lead us to a loss slower than other losing evaluations. The best losing evaluation would be a loss in ML moves. This makes intuitive sense when the opponent might not play perfectly. By trying to extent the game the chances of an imperfect player making a mistake increases which could change the game from losing to either a draw or winning situation.

From the information given so far about the evaluation space we observe the most important property of all the evaluations, namely that all evaluations are zero-sum. This means that an advantage for one player is an equivalent loss for the other player [7]. In other words, we have that any evaluation  $e_{\text{player}_1}$  for player 1 and  $e_{\text{player}_2}$  for player 2 at a certain node adds up to zero, i.e.

$$e_{\text{player}_1} + e_{\text{player}_2} = 0$$

<sup>1</sup>I have implemented everything in Python.

$$e_{\text{player}_1} = -e_{\text{player}_2}$$

This is extremely useful as we can now just negate an evaluation to change the perspective of the evaluation to the other player. This property is on the basis of negamax and often given as a requirement for minmax [3].

Now that we have taken a look at losses and wins and that we have discussed the zero-sum property it is time to take a look at draws. Draws are usually considered to be neutral in search programs and an equal outcome for both players. However, this does not reflect differences in player strength. For example a better player would rather not play a draw against a worse player, while for the worse player this would constitute a good result.

I have implemented negamax and various variants of negamax. Since negamax, when no heuristic is used, gives a perfect result it is safe to assume that the opponent will not be stronger. Because of this we can prefer later draws for the negamax player and earlier draws for the opponent. Considering this we give the positive draw space to the negamax player and the negative draw space to the opponent.

## 2.2 Heuristic

Apart from exact evaluations we can also have a heuristic evaluation. This represents an approximation of the outcome when to exact evaluation is unknown. In my program implementation I have an interface from which various heuristics can be defined. For every heuristic the zero-sum property must hold and the value must be between zero and a given limit. This limit is used in the evaluation space to create enough width for all the possible heuristic values.

In the evaluation space we have two gaps for heuristic values. One on the positive side and one on the negative side. This way we can negate the heuristic value to change to the other player's perspective.

Note that the heuristic space is outside of the draw space. This means that heuristic values which seem positive but can turn out into a loss are preferred over definite draws. If we have high confidence in our heuristic this is fine. If not we can also swap the heuristic and draw spaces. This is just a design choice that I have made.

## 3 Heuristics

My program contains two heuristics the given heuristic and a new heuristic.

### 3.1 Given Heuristic

Let me start by stating that the given heuristic does not satisfy the zero-sum property, and this is quite the problem. So the given heuristic provides the number of consecutive stones the current player on turn has. This however, does not say anything about the relative position compared to the other player.

Imagine that player 1 has three consecutive stones, then we have  $e_{\text{player}_1} = 3$ . Lets also say that we need four consecutive stones for a win, so  $HL + 1 = N = 4$ . There is no way that we can use this value to say anything about  $e_{\text{player}_2}$ . For example, we could say  $e_{\text{player}_2} = -e_{\text{player}_1} = -3$ . This would constitute the worst heuristic evaluation for player 2, however for all we know player 2 also has 3 in a row, therefore this is not a good translation. Another example could be to say  $e_{\text{player}_2} = HL - e_{\text{player}_1} = 3 - 3 = 0$ . This would be the best heuristic value for player 2, but maybe player 2 does not have any consecutive stones anywhere!

The problem with the given heuristic is that it only gives information about one of the two players. This makes it that we could never change this value in a meaningful way to represent a heuristic evaluation for the other player.

How would this be solved? In chess engines for example, a heuristic value often constitutes the *difference* in points given by the pieces on the board, where every piece is assigned a static value<sup>2</sup>.

<sup>2</sup>I am oversimplifying a lot here, see <https://www.chessprogramming.org/Score>.

If we would do this in Connect Four, we could take the maximum number of consecutive for both players and take the difference. This does satisfy the zero-sum property. For example player 1 has 3 consecutive stones and player 2 only 1, then the heuristic value for player 1 is  $e_{\text{player}_1} = 3 - 1 = 2$  and for player 2 is  $e_{\text{player}_2} = -e_{\text{player}_1} = -2$ . Now if player 2 gets more consecutive stones the heuristic value reflects the advantage and the loss for both player 1 and player 2.

I have implemented the heuristic as given in the assignment. This means that when this heuristic is used the outcome is incorrect. However, a fix would be trivial, namely by calling the heuristic twice and taking the difference. Albeit this would be extremely inefficient.

### 3.2 Nop Heuristic

I have also created a simple heuristic of my own so I can work with a depth bounded search. This heuristic always returns 0<sup>3</sup>. This way we prefer to go through a path which has an undefined outcome over a definite loss, but we prefer a definite draw<sup>4</sup> and definite win over an undefined outcome. This heuristic is meant as a filler to make depth bounded search work.

## 4 Negamax AI

The first algorithm that I have implemented is negamax, this is equivalent to minmax but more elegant. The negamax algorithm makes use of the fact that minmax applies to zero-sum games and the fact that  $\max(a, b) = -\min(-a, -b)$ . Making use of these properties we can get rid of the case distinction in the minmax algorithm. The negamax algorithm is shown in 1.

---

#### Algorithm 1 Negamax

---

**Require:** node, depth, color, timeout

**Ensure:** node evaluation

```

1: function NEGAMAX(node, depth, color, timeout)
2:   if timeout is reached then
3:     return undefined
4:   end if
5:   if node is over then
6:     return color  $\times$  evaluation of node
7:   end if
8:   if depth = 0 then
9:     return color  $\times$  heuristic evaluation of node
10:  end if
11:  value :=  $-\infty$ 
12:  for child of node do
13:    value := max(value,  $-\text{NEGAMAX}(\text{child}, \text{depth} - 1, -\text{color}, \text{timeout})$ )
14:  end for
15:  return value
16: end function

```

---

The algorithm works pretty straightforward. The color represents which person is on the move<sup>5</sup>. Note that I have added a timeout which when reached stops the execution of the algorithm, this is very useful when benchmarking or live interacting with the algorithm. The algorithm which calls the negamax function is a bit different and is shown in 2.

---

<sup>3</sup>In the code this is `Eval(const = Eval.UNDEFINED)`.

<sup>4</sup>Remember that the algorithm always has the positive draw space.

<sup>5</sup>In my implementation this is called `rootplayer`.

---

**Algorithm 2** Negamax move

---

**Require:** node, depth, timeout**Ensure:** best moves, node evaluation

```
1: function MOVE(node, depth, timeout)
2:   bestvalue :=  $-\infty$ 
3:   bestmoves :=  $\emptyset$ 
4:   for move at node do
5:     child := play move at node
6:     value := -NEGAMAX(child, depth - 1, -1, timeout)
7:     if timeout is reached then
8:       return undefined
9:     end if
10:    if value > bestvalue then
11:      bestvalue := value
12:      bestmoves := move
13:    else if value = bestvalue then
14:      bestmoves := bestmoves + move
15:    end if
16:  end for
17:  return bestmoves, bestvalue
18: end function
```

---

## 5 Alpha-beta AI

The second algorithm that I have implemented is alpha-beta pruning in the negamax algorithm. This is very similar to negamax, we just add an alpha-beta window. The alpha value represents the value that the current player is assured of, the beta value represents the value that the opponent is assured of. Whenever  $\alpha > \beta$  we know that the opponent would never let us go down this path as they already have a better move somewhere else, so we can prune that branch. With alpha-beta pruning we still get the same result as with pure negamax. Note that  $[\alpha, \beta]$  for player 1 is the same as  $[-\beta, -\alpha]$  for player 2, this is repeatedly used in negamax. Negamax with alpha-beta pruning is shown in 3, the move generation is shown in 4.

One peculiarity in my alpha-beta pruning implementation is that we do not want to have a cutoff when either the alpha or the beta value is undefined. This would lead to too early cutoffs and incorrect results.

## 6 Insane AI

The third algorithm that I have implemented is an extension of the alpha-beta pruning negamax algorithm with move ordering, a transposition table, iterative deepening and MTD(f).

### 6.1 Move ordering

The first change is move ordering. Optimally, the alpha-beta pruning algorithm is run with the best move first every time, this would cause the maximum number of cutoffs. Ofcourse, we do not know the best move at first, but an approximation can already greatly increase the speed of the algorithm. To this end I use a simply move ordering algorithm which prefers columns in the middle of the board over columns towards the edges of the board. This can be easily achieved by sorted the moves according to their value given by the formula  $w/2 - c$ , where  $w$  is the width of the board and  $c \in [0, w)$  the number of a column representing a move.

### 6.2 Transposition table

The second change is the addition of a transposition table, or TT for short. A TT is a memory construct which saves evaluations of states that we have visited. Whenever we visit a state a

---

**Algorithm 3** Negamax with alpha-beta pruning

---

**Require:** node, depth, color, timeout**Ensure:** node evaluation

```
1: function NEGAMAX(node, depth, color, alpha, beta, timeout)
2:   if timeout is reached then
3:     return undefined
4:   end if
5:   if node is over then
6:     return color  $\times$  evaluation of node
7:   end if
8:   if depth = 0 then
9:     return color  $\times$  heuristic evaluation of node
10:  end if
11:  value :=  $-\infty$ 
12:  for child of node do
13:    value := max(value, -NEGAMAX(child, depth - 1, -color, -beta, -alpha, timeout))
14:    if alpha  $\geq$  beta and alpha  $\neq$  undefined and beta  $\neq$  undefined then
15:      break
16:    end if
17:  end for
18:  return value
19: end function
```

---

---

**Algorithm 4** Negamax with alpha-beta pruning move

---

**Require:** node, depth, timeout**Ensure:** best moves, node evaluation

```
1: function MOVE(node, depth, timeout)
2:   alpha :=  $-\infty$ 
3:   beta :=  $\infty$ 
4:   bestvalue :=  $-\infty$ 
5:   bestmoves :=  $\emptyset$ 
6:   for move at node do
7:     child := play move at node
8:     value := -NEGAMAX(child, depth - 1, -1, -beta, -alpha, timeout)
9:     if timeout is reached then
10:      return undefined
11:     end if
12:     if value > bestvalue then
13:       bestvalue := value
14:       bestmoves := move
15:     else if value = bestvalue then
16:       bestmoves := bestmoves + move
17:     end if
18:   end for
19:   return bestmoves, bestvalue
20: end function
```

---

second time, we can take the evaluation from the TT and prune the branch.

The implementation of such a table has to be fast. To this end I made a primitive hashtable without collision resolution. This table is implemented as an array with a prime number sized length. It has two functions, one for storing a value and a flag and one for retrieving the value and flag, the pseudocode can be seen in 5.

---

**Algorithm 5** Table put and get function

---

```

1: function PUT(key, value, flag)
2:   index := key % table size
3:   table[index] = (key, value, flag)
4: end function
5: function GET(key)
6:   index := key % table size
7:   entry := table[index]
8:   if entry  $\neq$  null then
9:     if entry.key = key then
10:      return entry
11:   end if
12: end if
13: return null
14: end function

```

---

The key generation for the transposition table should also be incredibly fast. In order to achieve this I have made a representation of the board in a ternary number. This ternary number uniquely identifies the board. The number functions as a flattened two dimensional array. The digits in the ternary number represent the state of every cell, a 0 for empty, 1 for a stone of player 1 and 2 for a stone of player 2. For a move made in a column  $c \in [0, w)$  I can now update the key with  $key = key + (x + w \cdot y)^3 * (onturn + 1)$  with  $onturn \in [0, 1]$  and  $w$  the width of the board. This achieves a constant time key generation.

When we add the transposition table to the alpha-beta pruning negamax algorithm we have another complication. Some values may not be an exact evaluation because of alpha-beta cutoffs. Luckily for us these values can still provide a lowerbound or an upperbound. For example, when the best value we find in a negamax iteration is bigger than beta, we know a cutoff has occurred, and we know that the value that we found is a lowerbound on the actual value of the node. Similarly when the best value we find is smaller than the original alpha value we know that this is an upperbound for the actual value of the node.

Using all of the information we have about the transposition table we can now add this to our algorithm. The new algorithm is shown in 6 [5]. Note that the move generation algorithm is the same as in 4.

### 6.3 Iterative deepening

The third addition is iterative deepening. This combines the best of depth first search, low space complexity, with the best of breadth first search, optimality. Iterative deepening is relatively simple, the algorithm can be seen in 7. Using the iterative deepening algorithm is now as simple as replacing the call in algorithm 4 at line 8 with the iterative deepening function. Note that instead of stopping when we do not have an undefined value anymore, we can also check whether this is a heuristic value and continue until we find an exact value.

### 6.4 MTD(f)

At last I have added MTD(f) a type of zero-window search. This search iteratively executes negamax with a window of just one big, causing a lot of cutoffs. Depending on the return value we can conclude whether the actual value is below or above our initial guess. This way we narrow down our guess until we find the correct one. The algorithm for MTD(f) can be seen in 8 [4].

---

**Algorithm 6** Negamax with alpha-beta pruning, move ordering and a TT

---

**Require:** node, depth, color, timeout, table**Ensure:** node evaluation

```
1: function NEGAMAX(node, depth, color, alpha, beta, timeout)
2:   if timeout is reached then
3:     return undefined
4:   end if
5:   alpha_original := alpha
6:   entry := table.get(node.key)
7:   if entry  $\neq$  null then
8:     if entry.flag = lowerbound then
9:       alpha := max(alpha, entry.value)
10:    else if entry.flag = upperbound then
11:      beta := min(beta, entry.value)
12:    else
13:      return entry.value
14:    end if
15:  end if
16:  if node is over then
17:    return color  $\times$  evaluation of node
18:  end if
19:  if depth = 0 then
20:    return color  $\times$  heuristic evaluation of node
21:  end if
22:  value :=  $-\infty$ 
23:  moves := sorted moves from node
24:  for move in moves do
25:    child := move at node
26:    value := max(value,  $-\text{NEGAMAX}(\text{child}, \text{depth} - 1, -\text{color}, -\text{beta}, -\text{alpha}, \text{timeout})$ )
27:    if alpha  $\geq$  beta and alpha  $\neq$  undefined and beta  $\neq$  undefined then
28:      break
29:    end if
30:  end for
31:  if value  $\neq$  undefined then
32:    if value  $\leq$  alpha_original then
33:      flag := upperbound
34:    else if value  $\geq$  beta then
35:      flag := lowerbound
36:    else
37:      flag := exact
38:    end if
39:    table.put(node.key, value, flag)
40:  end if
41:  return value
42: end function
```

---



---

**Algorithm 7** Iterative deepening

---

**Require:** node, maxdepth, color, timeout, table

**Ensure:** node evaluation

```
1: function ITERDEEP(node, maxdepth, color, alpha, beta, timeout)
2:   value := undefined
3:   for depth in [0, maxdepth] do
4:     if timeout is reached then
5:       return value
6:     end if
7:     value := NEGAMAX(child, depth, -1, alpha, beta, timeout)
8:     if value  $\neq$  undefined then
9:       return value
10:    end if
11:  end for
12: end function
```

---

---

**Algorithm 8** MTD(f)

---

**Require:** node, depth, color, timeout, table

**Ensure:** node evaluation

```
1: function MTDf(node, depth, color, alpha, beta, timeout)
2:   lowerbound := alpha
3:   upperbound := beta
4:   guess := 0
5:   while lowerbound < upperbound do
6:     beta := max(g, lowerbound + 1)
7:     guess := NEGAMAX(child, depth, -1, beta - 1, beta, timeout)
8:     if guess < beta then
9:       upperbound := guess
10:    else
11:      lowerbound := guess
12:    end if
13:  end while
14: end function
```

---

## 7 Complexity

The negamax algorithm traverses the whole search tree. The size of the search tree is defined by the average branching factor  $b$  and the depth  $d$  of the tree. Unless the depth is bound by the user we have an upperbound of  $d = w \cdot h$  with  $w$  and  $h$  being the width and height of the board. An upperbound on the branching factor is  $b = w$ . This gives us a complexity of  $\mathcal{O}(b^d) = \mathcal{O}(w^{w \cdot h})$ .

The negamax algorithm with alpha-beta pruning provides an average improvement over the original negamax algorithm due to the beta cutoffs. In the worst case scenario there are no beta cutoffs and the complexity stays the same. In the best case scenario the move ordering for the search is optimal and the number of leaf nodes searched is about  $\Omega(b^{d/2})$  [1].

The negamax algorithm with further improvements does not impact the complexity significantly<sup>6</sup>. The move ordering makes the alpha-beta pruning more efficient and closer to the lower-bound given. The transposition table has constant lookup time, constant saving time and constant key generation. So the table does not introduce additional complexity, but it does cause more cutoffs. The complexity of iterative deepening is the same as the original negamax algorithm [2].

Regarding MTD(f), in the worst case scenario we increase our guess by one every time until we reach the actual value, we could do this a maximum of  $2 \cdot w \cdot h$  times, namely half the draw space and the win or losing space, assuming  $HL = 0$ . This gives us an upperbound of  $\mathcal{O}(w \cdot h \cdot w^{w \cdot h})$  which constitutes a minimal increase considering we use a transposition table and we get a large number of cutoffs. In the best case scenario we immediately guess the evaluation correctly and have the same as the original negamax search, but with a lot of cutoffs. It is extremely difficult to determine the average complexity of my improved algorithm as a lot of advanced statistical techniques are used in combination with each other, so this is out of scope of this project.

## 8 Benchmarks

Using the generator available in the code base I have made test cases which can be used to benchmark the different algorithms. The first number in the test set name denotes the number of moves already made. The second number in the test name is the maximum depth at which the solution can be found, with 7 having a minimum of 2 and 14 a minimum of 7.

All of the benchmarks are run with a timeout set at ten seconds. No depth bound is used and the default Connect Four settings are used. I did not have enough time to test for various Generalised Connect Four settings as generating test cases for this and the actual benchmarking takes a lot of time<sup>7</sup>.

### 8.1 Negamax AI

The benchmark results for the pure negamax AI can be seen in table 1. The count gives the number of nodes that have been visited. Interestingly for both test sets starting with 20 moves all of the test cases reached a timeout. This is because negamax has to search the whole tree which at that point is simply too big to do in just ten seconds.

test set	total	timeouts	mean time	max time	mean count	max count	count per sec
30-07	30	2	0.253901	9.589689	28040	1051029	110008
20-07	30	30	0	0	0	0	0
30-14	30	1	0.719904	8.746326	67918	832406	96602
20-14	30	30	0	0	0	0	0

Table 1: Negamax / minimax

<sup>6</sup>But it does greatly improve performance in practise!

<sup>7</sup>And I am already grossly over the word limit...

## 8.2 Alpha-beta AI

The benchmark results for the negamax AI with alpha-beta pruning can be seen in table 2. The count per second is slightly decreased compared to the plain negamax version, but this is probably because the search is now super short. The overall speed of the alpha-beta pruning version is way better, more than tenfold, compared to the plain version. This can be explained by the same tenfold decrease in nodes visited due to the cutoffs. However, we still get a lot of timeouts when the number of moves made is 20.

test set	total	timeouts	mean time	max time	mean count	max count	count per sec
30-07	30	0	0.031056	0.809639	2909	73283	90247
20-07	30	27	6.843755	9.205496	560175	824355	86068
30-14	30	0	0.058018	0.635897	4228	46229	79448
20-14	30	29	2.546612	2.546612	202736	202736	79610

Table 2: Negamax with alpha-beta pruning

## 8.3 Insane AI

For the insane AI I have run three benchmarks, one with iterative deepening, one with MTD(f) and one with both, see tables 3, 4 and 5. These results are actually very interesting.

We can see from these tables that iterative deepening in particular has a big influence on performance. Where we saw a lot of timeouts before at the test sets which start with 20 moves, we now actually have most of them solved within the ten seconds timeout. The lower the distance of the solution, the faster we will find it with iterative deepening. When the distance becomes larger the effect of iterative deepening will become smaller.

Interestingly, MTD(f) actually seems to create a decrease in performance. Both when used with and without iterative deepening. This is likely explained by the fact that the transposition table does not store all that many states.<sup>8</sup> Because of this the overhead of recomputing new zero-window searches is big.

test set	total	timeouts	mean time	max time	mean count	max count	count per sec
30-07	30	1	0.095417	2.011071	197	146619	38598
20-07	30	7	0.108103	2.660797	278	216104	52110
30-14	30	1	0.222108	5.05804	6886	323909	52655
20-14	30	4	0.921285	6.190709	61932	468116	67406

Table 3: Insane AI with iterative deepening

test set	total	timeouts	mean time	max time	mean count	max count	count per sec
30-07	30	0	0.13946	3.717605	1519	246017	45682
20-07	30	30	0	0	0	0	0
30-14	30	0	0.183881	0.819843	4072	50014	39368
20-14	30	28	8.018537	8.018537	670090	670090	81651

Table 4: Insane AI with MTD(f)

---

<sup>8</sup>This can be seen when using the interactive mode of the program, `maker`.

test set	total	timeouts	mean time	max time	mean count	max count	count per sec
30-07	30	1	0.095708	4.261424	377	339069	55753
20-07	30	8	0.118574	4.366139	660	335230	59456
30-14	30	1	0.415391	9.557776	20226	664495	61041
20-14	30	6	2.968389	9.57551	209221	643107	71115

Table 5: Insane AI with iterative deepening and MTD(f)

## 9 Conclusion

Comparing all of the implemented algorithms, we can conclude that alpha-beta pruning gives a big performance increase, as well as the move ordering, the transposition table<sup>9</sup> and iterative deepening, while MTD(f) does not.

## 10 Discussion

I have implemented everything in Python. This allowed me to do extremely fast prototyping. However, for search algorithms where speed is of utmost importance Python is everything but ideal. It would have been way better if we were allowed to use a low level language like C, C++ or Rust, this alone would greatly improve the capabilities of the AI's.

Since we had to implement all of the AI's for Generalised Connect Four it is way more difficult to implement a bitboard. This is a loss in efficiency, but also a gain in generalisability.

In the benchmark results we had a lot of timeouts. However, if I had more time I could have run the benchmarks with a depth bound. This would have given the algorithms a way better chance of finishing within the ten seconds timeout. Adding iterative deepening to the original negamax AI and the alpha-beta pruning version would have greatly increased their performance.

The usefulness of the transposition table could have been further increased by keeping a separate key for a mirror board. This is a relatively small addition, and could yield a significant performance increase, especially for searches where only a few moves are made and a cutoff is big.

Further heuristics could be added to the game. For example based on the board size  $w \cdot h$  you can determine whether we should aim for winning moves on even or odd rows.

## 11 References

- [1] Wikipedia contributors. Alpha-beta pruning — Wikipedia, the free encyclopedia. [https://en.wikipedia.org/w/index.php?title=Alpha%E2%80%93beta\\_pruning&oldid=1112997890](https://en.wikipedia.org/w/index.php?title=Alpha%E2%80%93beta_pruning&oldid=1112997890), 2022. [Online; accessed 7 October 2022].
- [2] Wikipedia contributors. Iterative deepening depth-first search — Wikipedia, the free encyclopedia. [https://en.wikipedia.org/w/index.php?title=Iterative\\_deepening\\_depth-first\\_search&oldid=1108524032](https://en.wikipedia.org/w/index.php?title=Iterative_deepening_depth-first_search&oldid=1108524032), 2022. [Online; accessed 7 October 2022].
- [3] Wikipedia contributors. Minimax — Wikipedia, the free encyclopedia. <https://en.wikipedia.org/w/index.php?title=Minimax&oldid=1113619084>, 2022. [Online; accessed 6 October 2022].
- [4] Wikipedia contributors. Mtd(f) — Wikipedia, the free encyclopedia. [https://en.wikipedia.org/w/index.php?title=MTD\(f\)&oldid=1112174230](https://en.wikipedia.org/w/index.php?title=MTD(f)&oldid=1112174230), 2022. [Online; accessed 6 October 2022].
- [5] Wikipedia contributors. Negamax — Wikipedia, the free encyclopedia. <https://en.wikipedia.org/w/index.php?title=Negamax&oldid=1067237200>, 2022. [Online; accessed 6 October 2022].

<sup>9</sup>The move ordering and TT performance increase are not tested in the benchmarks, but can be manually confirmed using the `maker`.

- [6] Wikipedia contributors. Solved game — Wikipedia, the free encyclopedia. [https://en.wikipedia.org/w/index.php?title=Solved\\_game&oldid=1112712449](https://en.wikipedia.org/w/index.php?title=Solved_game&oldid=1112712449), 2022. [Online; accessed 6 October 2022].
- [7] Wikipedia contributors. Zero-sum game — Wikipedia, the free encyclopedia. [https://en.wikipedia.org/w/index.php?title=Zero-sum\\_game&oldid=1105258284](https://en.wikipedia.org/w/index.php?title=Zero-sum_game&oldid=1105258284), 2022. [Online; accessed 6 October 2022].

## A Implementation

In this section I provide a few remarks on the implementation. Firstly, I have added an interactive environment to the program from which you can test all the algorithms and play around with all the settings. This interactive environment can be started by using the `maker`, use `?` to get a list of all the commands available.

Secondly, I have reimplemented the `isover` function so that it only checks if the game is finished due to the last move that has been made. Especially on bigger boards this can greatly increase the performance of this function as we do not have to check the whole board anymore.

Thirdly, I have not used an explicit tree structure but an implicit one which is present in the recursive structure of `negamax`. This is more space efficient. I was planning to also implement the Monte Carlo Tree search algorithm, and therefore also some kind of tree structure implementation for this, but due to time constraints I did not get to this.